

CS 6923 Machine Learning

Spring 2019

Final Project Report

Name: Kaustuk Kumar
NetID: kk3697

Name: Nikhil Narasimha Prasad
NetID: nnp267

PART I: Preprocessing (No more than two pages for this part)

1. How does your program handle missing values? And why?

- The features that have missing values are “house”, “weight”, “player_code”, “move_specialty”.
- Missing data can be treacherous because it is difficult to identify the problem. It is not always obvious when missing data will cause a problem.
- Missing data can also lead to misleading results by introducing bias. Whenever segments of the target population do not respond, they become under-represented in the data. In this situation, we end up not analyzing what we intended to measure.
- Weight is missing in over approximately 97% records. These missing values are poorly interpretable and cannot be generalized to other players. Therefore, this feature is just dropped completely.
- For the other features, we use an “Imputer” function which replaces all the missing values with the median of all the values in the feature.

```
from sklearn.preprocessing import Imputer

imp = Imputer(missing_values = 'NaN', strategy = 'median', axis = 0)
imp.fit(X_test_new)
X_test_new = pd.DataFrame(data = imp.transform(X_test_new), columns = X_test_new.columns)
```

2. If your program converts numeric features to categorical features, or categorical features to numeric features, describe how it does it.

- There are multiple categorical features in the data, which need to be converted to numeric features to run our classification models.
- The categorical features that need to be converted are “house”, “gender”, “player_code”, “move_specialty”, “player_type”, “snitchnip”, “stooging”, “change” and “snitch_caught”.
- We use cat.codes to assign numeric codes to these categorical features as show below.

```
X_new['house'] = X_new['house'].astype('category')
X_new['gender'] = X_new['gender'].astype('category')
X_new['player_code'] = X_new['player_code'].astype('category')
X_new['move_specialty'] = X_new['move_specialty'].astype('category')
X_new['player_type'] = X_new['player_type'].astype('category')
X_new['snitchnip'] = X_new['snitchnip'].astype('category')
X_new['stooging'] = X_new['stooging'].astype('category')
X_new['change'] = X_new['change'].astype('category')
X_new['snitch_caught'] = X_new['snitch_caught'].astype('category')

cat_columns = X_new.select_dtypes(['category']).columns
X_new[cat_columns] = X_new[cat_columns].apply(lambda x: x.cat.codes)
X_new.info()
```

3. Describe any feature selection, combination or creation, and any feature values combination performed by your program and the reasons for doing so.

- The features “num_games_satout”, “num_games_injured” and “num_games_notpartof” can be combined to create a new feature which indicates that the player did not participate.
- We create a new feature called “num_games_missed”, which is the sum of the above three features. After creating the new feature, we can drop the previous 3 features, hence simplifying our model.
- We calculate the total number of tactics used by every player and create a new feature called “total_tactics_used” from the 23 features for every tactic used by a player.
- We take the count of all the tactics features that have a value which is not equal to “No” and store the count in the new feature “total_tactics_used”.
- The feature “total_tactics_used” is then used for any future preprocessing or model fitting, and the 23 features for tactics are dropped from the dataset. This also helps in simplifying the model even further.

4. Describe other preprocessing used in your program (e.g. centralizing, normalization)

- We reduce the number of categories for the features “stooging” and “snitchnip”. These features are recategorized into 3 categories “normal”, “none” and “high”.
- We perform Log Transformation for the feature that is the most skewed in the dataset, i.e., “num_games_missed”.
- Since we used log transformation, we standardize the “num_games_missed” feature using the formula:
$$\text{New value} = (\text{Value} - \text{Mean}(\text{values})) / \text{StandardDeviation}(\text{values})$$
- We remove the outliers for all the numerical features which helps us in getting a more relevant mean and standard deviation.
- We restrict the data to 3 standard deviations on either side from the mean for each numeric value, as shown below.
- Since the data is very imbalanced, we use SMOTE to make the data balanced.

```
def find_outliers(x) :
    mean = np.mean(x)
    std = np.std(x)
    floor = mean - 3*std
    ceiling = mean + 3*std
    outlier_indices = set((x.index[(x < floor) | (x > ceiling)]))
    return outlier_indices

outlier_indices = find_outliers(X_new['age'])
outlier_indices = outlier_indices.union(find_outliers(X_new['foul_type_id']))
outlier_indices = outlier_indices.union(find_outliers(X_new['game_move_id']))
outlier_indices = outlier_indices.union(find_outliers(X_new['penalty_id']))
outlier_indices = outlier_indices.union(find_outliers(X_new['game_duration']))
outlier_indices = outlier_indices.union(find_outliers(X_new['num_game_moves']))
outlier_indices = outlier_indices.union(find_outliers(X_new['num_game_losses']))
outlier_indices = outlier_indices.union(find_outliers(X_new['num_practice_sessions']))
outlier_indices = outlier_indices.union(find_outliers(X_new['num_games_missed']))
outlier_indices = outlier_indices.union(find_outliers(X_new['foul_type_id']))
outlier_indices = outlier_indices.union(find_outliers(X_new['total_tactics_used']))
print(len(outlier_indices))
print(len(X_new))
print(len(y_new))
for i in outlier_indices:
    X_new.drop(i, inplace = True)
    y_new.drop(i, inplace = True)
```

PART II: Classification (No more than two pages for each model in this part)

Model One:

1. Supervised learning method used in this model is
Logistic Regression
2. Why you choose this supervised learning method?
 - Logistic regression is used to describe data and to explain the relationship between one dependent binary variable and one or more nominal, ordinal, interval or ratio-level independent variables.
 - We choose this model to estimate a baseline accuracy for the preprocessed dataset.
 - We assess the outputs of the data to make a decision if a more complex model is required.
3. Describe the method you used to evaluate this method.
 - Accuracy:
 - $\text{Accuracy} = (\text{TP} + \text{TN}) / (\text{TP} + \text{TN} + \text{FP} + \text{FN})$
 - We use the `accuracy_score()` method provided by sk-learn to get the accuracy of the model on the the given data.
 - Confusion matrix
 - A confusion matrix is a table that is often used to describe the performance of a classification model on a set of test data for which the true values are known.
 - We use the `confusion_matrix()` method provided by sk-learn to get the confusion matrix of the model on the given data.
4. Describe process of experimenting different parameter settings or associated techniques.
 - Parameter name: `class_weight` (Weights associated with classes in the form `{class_label: weight}`)
 - Parameter values: None, balanced
 - Performance of different values:
 - `Class_weight = None`, Accuracy = 48.28%
 - `Class_weight = balanced`, Accuracy = 48.28%
 - Analysis:
 - The performance is similar for both the values in the case of logistic regression.

- **Parameter name: C (Inverse of regularization strength)**
 - **Parameter values: 0.01, 1, 100**
 - **Performance of different values:**
 - **C = 0.01, Accuracy = 48.28%**
 - **C = 1, Accuracy = 48.28%**
 - **C = 100, Accuracy = 48.28%**
 - **Analysis:**
 - **The performance is similar for both the values in the case of logistic regression.**
- **Parameter name: folds (Number of folds in K-Fold Cross Validation)**
 - **Parameter values: 5, 10, 15, 20**
 - **Performance of different values:**
 - **folds = 5, Accuracy = 45.43%**
 - **folds = 10, Accuracy = 46.18%**
 - **folds = 15, Accuracy = 47.46%**
 - **folds = 20, Accuracy = 48.28%**
 - **Analysis:**
 - **The performance of this model improves as the number of folds increases.**

5. Accuracy and Confusion matrix with most suitable parameters

		Predicted	
		-	+
Actual	-	44143	38205
	+	46973	35375

Accuracy: 48.28%

Model Two:

1. Supervised learning method used in this model is
Decision Tree Classifier
2. Why you choose this supervised learning method?
 - The general motive of using Decision Tree is to create a training model which can use to predict class or value of target variables by learning decision rules inferred from training data.
 - Decision trees implicitly perform variable screening or feature selection.
 - Decision trees require relatively little effort from users for data preparation.
 - Nonlinear relationships between parameters do not affect tree performance.
3. Describe the method you used to evaluate this method.
 - Accuracy
 - $\text{Accuracy} = (\text{TP} + \text{TN}) / (\text{TP} + \text{TN} + \text{FP} + \text{FN})$
 - We use the `accuracy_score()` method provided by sk-learn to get the accuracy of the model on the the given data.
 - Confusion matrix
 - A confusion matrix is a table that is often used to describe the performance of a classification model on a set of test data for which the true values are known.
 - We use the `confusion_matrix()` method provided by sk-learn to get the confusion matrix of the model on the given data.
4. Describe process of experimenting different parameter settings or associated techniques.
 - Parameter name: leaves (Number of leaves in the decision tree)
 - Parameter values: 10, 30, 100, 500
 - Performance of different values:
 - leaves = 10, Accuracy = 91.43%
 - leaves = 30, Accuracy = 91.82%
 - leaves = 100, Accuracy = 91.37%
 - leaves = 500, Accuracy = 88.98%

- **Analysis:**
 - Initially the performance of the model improves with increase in number of leaves, peaks at 30, and then the accuracy decreases with increase in number of leaves.
- **Parameter name: min_samples_split** (Minimum number of samples required to split an internal node)
 - **Parameter values:** 2, 5, 10
 - **Performance of different values:**
 - min_samples_split = 2, Accuracy = 91.82%
 - min_samples_split = 5, Accuracy = 91.82%
 - min_samples_split = 10, Accuracy = 91.82 %
 - **Analysis:**
 - The performance is similar for all values of min_samples_split
- **Parameter name: max_depth** (The maximum depth of the tree)
 - **Parameter values:** 10, 30, 100
 - **Performance of different values:**
 - max_depth = 10, Accuracy = 89.58%
 - max_depth= 30 Accuracy = 91.82%
 - max_depth = 100, Accuracy = 91.82 %
 - **Analysis:**
 - Initially the accuracy increases with increase in depth of the tree, but stagnates after a depth of 30.

5. Accuracy and Confusion matrix with most suitable parameters

		Predicted	
		-	+
Actual	-	16084	302
	+	2391	14163

Accuracy: 91.82%

Model Three:

1. Supervised learning method used in this model is
Random Forest Classifier
2. Why you choose this supervised learning method?
 - Random Forest builds multiple trees and combines them together to get a more accurate result.
 - Random forest classifier will handle the missing values and maintain the accuracy of a large proportion of data.
 - If there are more trees, it won't allow overfitting trees in the model.
 - It has the power to handle a large data set with higher dimensionality
3. Describe the method you used to evaluate this method.
 - Accuracy
 - $\text{Accuracy} = (\text{TP} + \text{TN}) / (\text{TP} + \text{TN} + \text{FP} + \text{FN})$
 - We use the `accuracy_score()` method provided by sk-learn to get the accuracy of the model on the the given data.
 - Confusion matrix
 - A confusion matrix is a table that is often used to describe the performance of a classification model on a set of test data for which the true values are known.
 - We use the `confusion_matrix()` method provided by sk-learn to get the confusion matrix of the model on the given data.
4. Describe process of experimenting different parameter settings or associated techniques.
 - Parameter name: `n_estimators` (The number of trees in the forest)
 - Parameter values: 100, 200, 300
 - Performance of different values:
 - `n_estimators = 100`, Accuracy = 94.00%
 - `n_estimators = 200`, Accuracy = 93.95%
 - `n_estimators = 300`, Accuracy = 94.05%
 - Analysis:
 - There is a minimal increase in the accuracy as the number of trees increases.

- **Parameter name: max_depth** (The maximum depth of the tree)
 - **Parameter values:** 30, 70, 90
 - **Performance of different values:**
 - max_depth = 30, Accuracy = 94.00%
 - max_depth = 70 Accuracy = 94.05%
 - max_depth = 90, Accuracy = 94.05%
 - **Analysis:**
 - There is a minimal increase in the accuracy as the max_depth increases and then in stagnates after 70.
- **Parameter name: min_samples_leaf** (The minimum number of samples required to be at a leaf node)
 - **Parameter values:** 1, 2, 4
 - **Performance of different values:**
 - min_samples_leaf = 1, Accuracy = 94.05%
 - min_samples_leaf = 2 Accuracy = 94.00%
 - min_samples_leaf = 4, Accuracy = 93.95%
 - **Analysis:**
 - The accuracy of the model decreases as the number of min_samples_leaf increases.

5. Accuracy and Confusion matrix with most suitable parameters

		Predicted	
		-	+
Actual	-	16346	40
	+	1919	14635

Accuracy: 94.05%

PART III: Best Hypothesis (No more than two pages for each model in this part)

1. Which model do you choose as a final method?

Model number: Three

Supervised learning method used in this model: Random Forest Classifier

2. Reasons for choosing this model.

- **This model gives us the best score for accuracy on the training data.**
- **This model also gives us the best confusion matrix amongst all the other models. It has the least number of False Positives and False Negatives.**

3. What are the reasons do you think that make it has the best performance?

- **Random forests are a strong modeling technique and much more robust than a single decision tree or logistic regression.**
- **Random forests aggregate many decision trees to limit overfitting as well as error due to bias and therefore yield better results.**